

Nexus Loop

A deterministic detect, diagnose, prescribe, verify loop for AI agent deployments

Team Name: Mayush Aura
Institution: Indian Institute of Technology, Patna

Abstract

We present **Nexus Loop**, a deterministic pipeline that turns execution logs from a deployed AI support agent into a small set of verified improvement decisions. The system follows the **Detect** → **Diagnose** → **Prescribe** → **Verify** workflow. It measures how much of the traffic each metric can cover, detects genuine regressions, dismisses lookalike changes, quantifies impact with explicit fidelity and coverage, links a cause to the responsible configuration change, proposes one fix per regression, evaluates the strongest fix on the replay service, and records the final human decision. On the practice corpus it scores **55.0 of 55** machine points, and the same on every held-out corpus we generated, with **about 9.2 seconds** runtime and per-metric coverage of **0.72 and 0.93** on tool metrics and **1.0** on cost and knowledge-base metrics. No language model is called for any fact the logs already record. We additionally report the measurement gaps we refused to fill, the calibration of the corpus judge against human labels, the defects we found in our own pipeline, and the limits of the system at larger scale.

Keywords: Agent Observability, Root Cause Analysis, AI Agents, Anomaly Detection, Evaluation, Human-in-the-Loop

1 Problem and Objective

A deployed agent writes a lot of data: conversations, tool calls, retrieval lookups, model calls, end states, and a timeline of configuration changes. The task is to turn that into a closed improvement loop rather than a manual investigation.

Our system addresses four stages:

1. **Detect:** identify genuine problems and affected traffic slices.
2. **Diagnose:** determine why the problem occurred and identify the relevant configuration change.
3. **Prescribe:** propose one concrete fix and state its associated risk.
4. **Verify:** replay the affected cohort and determine whether the fix actually improves the outcome.

An operator reading the final screen should be able to answer six questions: what broke, who was affected, why it happened, who should act, what happens if nobody acts, and whether the proposed fix can be trusted.

1.1 Objectives

- Detect genuine regressions from deployment logs.
- Identify the precise traffic slice affected.
- Quantify impact using a denominator we can defend.
- Diagnose the underlying cause and link it to a configuration change.
- Explicitly dismiss lookalike and false-alarm patterns with a written reason.
- Refuse questions that cannot be answered from the available data.
- Mine the deployment’s own good behaviour as a standard.
- Propose a fix and verify it through the replay service.
- Keep a human approval or rejection decision in the loop.

2 Data and System Understanding

2.1 Available Data

The corpus covers 56 days across two tenants, acme-bank and northwind-retail. It contains 80,252 sessions and 883,764 execution steps. The step mix is 368,517 model calls, 365,065 turns, 72,168 knowledge-base lookups, 62,198 tool calls, 11,854 handoffs, and 3,962 guardrail events. There are 16 timestamped configuration changes across the two tenants, 400 outcome labels, and 96 human rubric labels. Work started on a smaller whole-session sample and the reported findings are confirmed on the full corpus.

2.2 Data Model and Joins

The data has three grains. A session is one whole conversation. A step is one internal action. A turn is one text exchange. Steps fan out many to one onto sessions, so any join that multiplies steps before an outcome is computed will double count the busier sessions. Our loader folds every step fact down to one row per session before any rate is computed, which removes that class of error at the source. `session_id` is the only join key we use, and `custom_dims.customer_ref` is close to unique per session and is not a key.

2.3 Measurement Constraints

Not every number can be measured from the logs, and the report marks this on every metric. Each metric carries a fidelity tag, a coverage figure, and, where the value is a judgement rather than a measurement, a calibration.

- **Fidelity:** measured, derived, or judged. Measured means we read or counted it from the logs. Derived means we computed it from measured primitives. Judged means a human or a judge decided meaning.
- **Coverage:** the share of traffic that can produce the metric. The legacy v2 flow emits no tool calls, so tool metrics cover 0.72 of acme-bank and 0.93 of northwind-retail traffic. Sessions on an agent kind that never emits the step type leave the denominator instead of being counted as zero.
- **Calibration:** for judged values, the agreement between the corpus judge and human labels, published with the sample size and per judge version.

Two questions are answered the other way on purpose. A tool failure rate is measured because we executed the call. The reason a user abandoned a session is a judgement about meaning, so it is filed as a gap rather than inferred from an end code.

2.4 Declared versus Measured Coverage

The catalog is a planning contract, and in two places it does not match the corpus we were given. We measure coverage from the corpus and publish the difference instead of trusting the prose.

This matters for honesty rather than for the score. A metric whose coverage is copied from the catalog would either overstate the hole or understate it. Measuring from the corpus keeps the denominator right.

3 System Overview

3.1 Architecture

The pipeline is one Python process. A kit loader streams the gzipped session and step files and keeps unknown fields. A feature layer folds steps to session grain. A coverage layer records which agent kinds emit which step types. Three fault detectors and three lookalike detectors inspect the session table. An asks layer answers the eleven questions as metrics or structured gaps. A standards layer mines a good-behaviour standard and freezes a golden set. The report layer assembles findings, diagnoses and prescriptions, validates the result against the provided schema, and writes it atomically. A replay client verifies one fix. A small standard-library server renders the report on a decision screen and writes the operator’s verdict back into the report.

3.2 End-to-End Pipeline

1. Data ingestion and validation
2. Session-grain feature table
3. Per-tenant emission coverage
4. Three fault detectors

Table 1: Important data entities and their usage.

Entity	Fields we use	Purpose
Session	tenant, day, intent, agent_id, agent_kind, session_end, hand-off_by_design, turns, cost_usd, quality_score, judge_version, milestones_reached	Outcome, cost, quality and cohort definitions
Step	step_type and per-type fields (tool status, result_field_count, retry_count, duration; kb_hit, kb_top_score; llm outcome)	Mechanism evidence, folded to session grain
Config timeline	day, tenant, kind, target, from_value, to_value	Attribution of a regression to a deploy
Catalog	fields, capabilities, cardinality budgets, milestones	Coverage planning and journey definitions
Human labels	session_id, human_quality	Calibration of the corpus judge

Table 2: Catalog claims checked against the corpus.

Claim in the catalog	What the corpus shows	How we handled it
Cost coverage 0.72 and 0.93, because v2 lacks cost	v2 sessions carry cost_usd and emit kb and guardrail steps	Measured coverage is 1.0; the metric carries observed_vs_declared
Calibration set of 100 human-scored sessions	manifest and labels are 96	Calibration published on the 96 labels that join
v2_flow emits no tool_call	Only the tool_call hole is real	This claim holds; tool coverage is 0.72 and 0.93

5. Three lookalike detectors
6. Eleven asks as metrics or gaps
7. Golden set frozen before any fix exists
8. Report assembly and schema validation
9. Replay verification of the strongest fix
10. Decision screen
11. Human decision written back into the report

The ordering matters in one place. The detectors run before the asks, so the asks can mark their windows as confounded when a detected event overlaps them.

4 Detection

4.1 Metrics

4.2 Detection Strategy

Each detector owns its own cohort dimensions. The tool detector groups by tenant and tool and then narrows to the dominant intent when one intent accounts for at least half of the affected sessions. The knowledge-base detector groups by tenant and intent. The efficiency detector groups by tenant and agent. No practice value such as a tenant name, a day, a tool or an agent appears in the detection code.

A shared window scanner finds the day window whose inside mean differs most from the outside mean, scored with a Welch z. This suits faults that start and end, and an open-ended window also catches a fault that has not recovered. The scan requires a minimum number of sessions inside and outside and at least a week of contrast outside, and windows open and close on days that carry traffic.

Thresholds are relative or statistical rather than tuned to this corpus. The tool detector requires an in-window empty rate of at least 5 percent, an absolute lift of at least 4 points, a relative lift of at least three times, and a z of at least 4. The resolution detector requires a drop of at least 5 points and a z of at least 4, with a stricter bar for a regression it cannot explain. The efficiency detector requires a turns ratio of at least 1.2 and a z of at least 5.

Two things about the scan bound how much weight a z deserves. The scan maximises over a large number of candidate windows and cohorts, so an unadjusted z of 4 is not a one-in-thirty-thousand event across the whole search. Sessions also cluster by day, agent and version, so treating them as independent draws makes the z optimistic. The high bars and the absolute-lift requirements are our guard against both, and we treat the effect size as the primary signal.

4.3 Affected Slice

Each finding names the slice it applies to. The tool fault is scoped to northwind-retail, intent order_status, tool get_order_status. The knowledge-base fault is scoped to acme-bank, intent premium_card_info. The efficiency fault is scoped to acme-bank, agent acme_main_v3. The lookalike findings are scoped the same way, and the judge-version change is scoped to both tenants.

5 Findings

5.1 Real Problems

For F-02 the impact block reads: 2,686 conversations, 29.0 percent of the tenant’s traffic in the window, 12 days, 94.31 dollars of recorded cost, 213 conversations that would have resolved at baseline, 404 unplanned handoffs, and 147 abandonments. For F-03 it reads 1,224 conversations, 23.3 percent of traffic, 10 days, 97.85 dollars, 2,448 extra turns, 983 slowed resolutions and 43.92 dollars of excess spend. Every finding names the owner who should act and the consequence of waiting. The owners differ because the causes differ: F-01 sits with the agent builder and the business owner who own the knowledge base, F-02 with the agent builder and the platform owner who own the tool contract, and F-03 with the agent builder who owns the prompt.

5.2 False Alarms and Lookalikes

Three changes look like regressions and are not. Each is reported as a non-regression finding with a linked diagnosis and a written reason, so a reader can see that we examined it.

Specificity is earned by looking and being right rather than by staying quiet, so these three dismissals are a first-class part of the result.

6 Diagnosis

6.1 Root-Cause Analysis

For each regression the diagnosis states the observed change, the cohort, the window, the candidate causes, the evidence that selects one, and the configuration change where one applies. The chain is recorded in the report and shown on the screen.

For F-02 the chain is: cohort resolution fell to 0.7755; the affected slice is order_status on get_order_status; the mechanism is an empty success, a 200 response with an outcome of ok and no result fields, on 13.7 percent of calls in the window, with a same-tool retry appearing in the affected sessions; the ordinary error rate stayed flat, which is why the fault was invisible; and the onset aligns with the get_order_status deploy from 3.2.0 to 3.3.0 on day 40.

Observation and attribution are separate. A regression is reported when the behaviour proves it, even when no configuration change aligns. In that case the cause is left unknown and no change is blamed.

6.2 Baseline Selection

The baseline is chosen per finding and stated in words on the impact block, so a reader can check it. F-02 and F-03 use the cohort’s own history over the 14 days before the window. F-01 has no before-period because the intent first appears on day 35, so it uses the other knowledge-base backed intents on the same agents in the same tenant and the same days, and its window opens on the launch change that created the cohort rather than on the first day that cleared the sample-size floor. Opening on the launch removes most of the detection lag. Where a cohort is new, we say so in the evidence and point to the peer baseline.

Nexus Loop — deterministic pipeline (stdlib only; no model for logged facts)




Figure 1: Overall architecture of the Nexus Loop system.

Table 3: Primary metrics used by the detection system.

Metric	Definition	Fidelity / Coverage
Containment rate	Share of sessions that resolved or handed off by design	Measured, coverage 1.0
Resolution rate	Share of sessions that ended resolved	Measured, coverage 1.0
Empty tool success	Tool calls with status 200, outcome ok, and result_field_count 0	Derived, coverage 0.72 and 0.93
Tool failure rate	Tool calls with outcome other than ok	Measured, coverage 0.72 and 0.93
KB confident-hit rate	Lookups that returned a document above the confidence bar	Measured, coverage 1.0 here
Turns to resolve	Median turns among resolved sessions	Measured, coverage 1.0
Cost per resolved conversation	Intent spend divided by its resolutions	Measured, coverage 1.0 here
Judged quality	Rubric score, segmented by judge version	Judged, calibrated, coverage 1.0

6.3 Configuration Change Attribution

Attribution reads which configuration dimension actually changed between the baseline and the window, then names the nearest matching change. This avoids blaming an innocent deploy that happens to share the onset day.

Table 6: Attributed configuration changes.

Finding	Kind	Target	From → To	Day
F-01	kb	premium_card_info	kb_2026_06_a → kb_2026_06_c	34
F-02	tool	get_order_status	3.2.0 → 3.3.0	40
F-03	prompt	acme_main_v3	p_v3 → p_v4	46

The attribution window is three days before the onset and two after, with the earlier change winning a tie. When the mechanism is proven but nothing aligns, the cause stays unknown and the change is left null.

7 Measurement Gaps

7.1 Unanswerable Questions

Two of the eleven questions have no honest numeric answer in this deployment, and one is a cohort that the planner should refuse.

We report the measured abandonment rate, 5.92 percent at acme-bank and 5.73 percent at northwind-retail, inside the gap text rather than as a metric. A measured metric tagged A09 would misread the question, which asks about the reason, not the rate.

7.2 How We Refuse

A refusal names what is missing, the event that would enable the measurement later, the near proxy an impatient team might use, and why that proxy misleads. For A11 the proxy is `retry_count > 0`, 1,956 rows in this corpus. Those are quality retries of the same target, not a switch to an alternate target, so reporting them as failover would show a small non-zero rate that reads as good news when failover does not exist. For A09 the proxy is a high turn count or a same-tool retry, which flags effort rather than frustration: a happy exit after two turns and a rage quit after two turns look the same in the logs. The same discipline applies when a judged metric cannot be calibrated. If fewer than ten human labels join, the judged quality reading is

replaced by an A06 gap that asks for a calibration set, rather than publishing a number with no agreement behind it.

8 Deployment Standard

8.1 Mining Good Behaviour

The standard comes from the deployment’s own sessions rather than a fixed target. A session is eligible when it resolved, had no unplanned handoff, and, where the telemetry exists for its agent kind, had no tool error, no empty tool answer and no knowledge-base miss. From the eligible baseline sessions we take the top decile by efficiency, or the 90th percentile of daily rates for resolution, and report the median as well. The result is an aspirational bar, not a typical one, and the derivation string says so.

The golden set is 30 eligible baseline sessions per finding, frozen to disk before any prescription exists and reused verbatim on later runs. Freezing before the fix removes hindsight: we cannot choose the known-good sessions after seeing which fix helps. One caveat we state plainly is that the golden set is the lowest-turn eligible tail, so passing it is a floor on quality rather than proof of broad improvement.

8.2 Deviation from Standard

Each affected cohort is compared with the standard for the same metric and cohort. F-02 resolved at 0.7755 against a standard median of 0.8545 and a best of 0.8926. F-01 resolved at 0.3087 against a standard median of 0.7273 and a best of 0.7399. F-03 took a median of 6.0 turns against a standard median of 4.0 and a best of 2.0.

9 Prescription

9.1 Proposed Fix

Each regression gets one concrete intervention, framed as a decision a human approves.

9.2 Risk and Shipping Criteria

Each prescription states what is being approved, the risk if the diagnosis is wrong, and the condition under which we would not ship. P-01 is low risk in the sense that failing to fix it changes nothing, but it carries a real risk the other way: if an empty body is a valid

Detector evidence — each series vs the window the scan selects (shaded)




Figure 2: Detector evidence: each daily series against the window the scan selects (shaded).

Table 4: Summary of genuine findings.

ID	Affected Slice	Finding / Impact
F-01	acme-bank, premium_card_info, days 34 to 44	Knowledge gap. Resolution fell to 0.3087 against a peer baseline of 0.7073. The cohort had no traffic before day 35, so its own history could not serve as a baseline. KB hit rate was 0.000 against 0.899 on peer intents.
F-02	northwind-retail, order_status, get_order_status, days 40 to 51	Silent tool failure. 13.7 percent of successful calls returned an empty body against 0.0 percent outside the window. Resolution fell from 0.8547 to 0.7755. The declared tool error rate stayed flat, so alerting on it would have missed the fault.
F-03	acme-bank, acme_main_v3, days 46 to 55	Prompt regression. Median turns to resolve rose from 4.0 to 6.0 and cost per session from 0.04406 to 0.07994 dollars, while resolution stayed flat at 0.7943 to 0.8031.

state for a brand-new or cancelled order, the fix would route valid sessions to a human. P-02 is low risk because added correct content does not remove anything. P-03 is moderate because the confirmation wording was added for safety, and if it is a compliance requirement then removing it trades cost for risk, which is a business decision.

Each prescription also carries a recovery estimate that closes nine tenths of the deficit, leaving the rest to genuine edge cases. The system changes the agent, and it never changes the metric definition, the judge version, or the reviewed good set.

10 Verification

10.1 Replay Method

The replay endpoint verifies a hypothesis we already hold, so we diagnose offline first and send one high-confidence run rather than probing. The request carries the exact detected cohort, the proposed change, and the frozen golden set. The endpoint can return an improvement, no effect, or a golden-set regression, and all three are recorded. A no-effect result is a result, not an error. If the service is unreachable the report stays valid with no verifications.

10.2 Before and After

Table 9: Replay verification results, P-01.

Metric	Before	After
Resolution rate	0.7741	0.8574
Golden set, 30 sessions	pass	pass
Prediction error	predicted +0.0713	replayed +0.0833

The replay returned improved, with a change of +0.0833 and a prediction error of -0.012 , and 0 of 30 golden-set sessions regressed. The before figure is the endpoint’s own simulated baseline, not our corpus measurement, and the report says so in the verification note. The prediction error compares our predicted delta with the replayed delta, and it is published on every verification.

10.3 Verification Decision

P-01 improved and is a candidate for deployment, with the golden set intact. The other two prescriptions carry no replay this run; the report records that they remain decision records without attached verification.

11 Decision Layer

11.1 Approve and Reject Workflow

The screen renders only the report. A reviewer sees the claim, the harm numbers, the cause, the proposed change, the replay result, and the risk, then records a verdict. Approve, reject and defer are all first-class. A rejection with a strong reason is a valid outcome, not a failed one.

The verdict writes into the report. The server serializes decisions on a lock, updates only the target prescription’s approval field, revalidates the whole report, writes it through one atomic helper, and appends the decision to a log. A reason is required.

11.2 Decision Record

Decisions live in the report itself. Each verdict, its author, its reason and a UTC timestamp are written into the corresponding prescription’s approval field, and every decision is appended to an append-only decisions.log beside the report. The screen is the only writer, and it revalidates the report before it commits.

12 Self-Assessment

12.1 Detection Performance

All three genuine faults were detected at the onset day, with the right cohort and cause, and no false alarms. Detection latency is zero because each window opens on the onset rather than on the first day that clears the sample floor.

Table 5: Lookalikes examined and rejected.

Case	Observed Change	Reason Dismissed
L-01	acme-bank aggregate containment moved +1.9 points on days 22 to 31	branch_locator’s share of traffic rose from 13 to 44 percent while every per-intent containment rate stayed flat within 5 points. This is a mix shift.
L-02	northwind-retail latency and volume spiked on days 12 to 15	Volume ran at 3.5 times normal and tool p95 latency rose from 1,338 to 3,152 ms, while resolution stayed flat. This is a load event for the platform owner.
L-03	Judged quality fell from 3.89 to 3.29 on day 28 across both tenants	The judge version changed from v1 to v2 on that day. We measured the rubric, not the agent, so we never trend quality across a judge boundary.

Diagnosis chain — F2 silent tool contract break (evidence at every step)




Figure 3: Diagnosis chain from observed regression to responsible configuration change, with evidence at each step.

12.2 Judge Calibration

No model is trained or used for judgement. The judged quality values come from the corpus’s own judge, and we publish the agreement between that judge and the human rubric labels.

Table 10: Judged quality calibration.

Item	Value
Judge versions	v1 and v2
Human-labelled sessions that join	96
Agreement, $ \Delta \leq 1$ on the 1 to 5 scale	0.9479
Mean absolute error	published in the report

A calibration of 1.00 is treated as a defect rather than a result, because the human labels disagree with one another. If agreement reaches 0.995 we publish the tighter half-point reading or withhold the calibration instead.

12.3 End-to-End Loop Completeness

Table 11: Automation coverage of the improvement loop.

Stage	Automated?	Implementation
Detect	Yes	Three detectors over the session table
Diagnose	Yes	Mechanism evidence plus the changed configuration dimension
Prescribe	Yes	One decision-shaped fix per regression
Verify	Yes	One high-confidence replay with the frozen golden set
Decision	Human	Approve, reject or defer on the screen, written back into the report

12.4 Threats to Inference

The detectors use statistics, so we state where those statistics are weaker than they look.

13 Implementation

13.1 Technology Stack

Table 13: Implementation stack.

Component	Implementation
Language	Python 3, standard library only
Data processing	Streaming gzip JSONL and CSV; one in-memory session table
Detection	Day-window scan with a Welch z, plus relative thresholds
Diagnosis	Mechanism evidence plus the changed configuration dimension
Semantic judgement	None. No model is called for any logged fact
Replay	HTTP client to the provided replay service
Screen	Standard-library HTTP server plus one HTML page
Report generation	Deterministic JSON assembly and schema validation
Storage	The report file, a decision log, a replay log, and a frozen golden set

13.2 Report Generation

The report is produced by the pipeline, validated against the provided schema plus a set of cross-field checks, and written through a unique temporary file with an atomic replace. The screen never recomputes a number. A decision updates one approval field and revalidates before it writes. Recorded approvals survive a regeneration: when the pipeline runs again it carries an approval forward onto a prescription with the same id, diagnosis and target, so a re-run does not silently erase a human decision.

13.3 Execution Environment

Python 3.9 or later, no third-party packages, and no network access except the optional local replay call. The pipeline runs offline. The screen is the only component that may reach the network, and only for optional web fonts; offline it falls back to system fonts, and it sends no data out.

14 Computational Efficiency

All figures below are the practice corpus (80,252 sessions, 883,764 steps). The sealed corpus has the same shape.

Table 7: Measurement gaps.

Question	Why Unanswerable	Required Event
A11 failover rate	No failover mechanism exists, so no failover event is ever emitted. There is no primary and secondary path to observe.	failover at step grain, fields from_target, to_target, reason, recovered, owner conversation-runtime
A09 abandonment reason	The abandonment end state is measured, but the reason is a judgement about meaning that no calibrated judge in the deployment answers.	abandonment_reason_judgement at session grain, fields judge_version, reason_class, confidence, calibration_set_id, owner quality-judging
A02 per-customer breakdown	custom_dims.customer_ref has 80,252 distinct values against a declared budget of 200. Grouping by it is a planner error.	custom_dims.segment, a closed set with a budget of 50, is the correct cohort key




Figure 4: Comparison between the mined deployment standard and the affected cohorts, split by metric so the units match.

Table 14: Computational performance.

Metric	Value
Average runtime	9.2 s, over three runs
Best runtime	9.1 s
Worst runtime	9.3 s
Peak memory	about 488 MB
Dataset pre-processing (load and fold)	about 5.3 s
Map / index / session table dominates; golden set a few kilobytes cache size	

The replay call adds nothing measurable at this size, one localhost request per run. The dominant cost is the single pass over the 883,764 steps and the session map it builds, about 5.3 s of the 9.2. The asks pass is next at about 2.3 s, mostly the calibration join and the review-queue scan. The three detectors together are about 1.3 s, so the window scan is not the bottleneck here even though it grows with the square of the number of days. A larger corpus changes that balance, because the session map is held in memory and the scan runs across every cohort.

15 Final Results

Table 15: Final system performance.

Metric	Result
Diagnostic accuracy	20 / 20
Specificity, false alarms	15 / 15, zero false alarms
Loop completeness	8 / 8
Honesty	12 / 12
Machine subtotal	55.0 / 55
Held-out pool	3 / 3 pass, 55.0 each
Moved-fault corpus	55.0
Moved-fault with a third tenant	55.0

All numbers are the practice corpus and the held-out corpora we generated; the sealed corpus is executed by the organisers against hidden ground truth. The held-out corpora were generated by a separate harness that renames labels, relocates the faults across tenants and slices, and adds an unseen tenant. Passing them with the same code, and with no corpus literal in the detectors, is the evidence that the result comes from the method rather than from a fit to the practice data.

16 Operator Screen

The screen is one brief and one signature rather than a dashboard. The header carries the coverage and fidelity chips and a disclosure for what is real, stubbed and simulated. A one-line switcher has three

views: **To decide** (the regressions), **Not a problem** (the lookalikes we examined and dismissed), and **We will not invent** (the refusals, failover first). The brief shows the claim, three harm numbers, the consequence of waiting, the cause, the proposed change and the replay. Evidence and baseline working sit under a disclosure, and the detector codes sit in a Technical names footer. The sticky panel carries the ask, the risk and the no-ship condition, and a traffic-light gate: a green Accept, a red Reject, and a grey Defer, each with a required reason. Arrow keys and Home/End move the selection through a long queue. The claim is set in Newsreader and the chrome in Source Sans 3, with system fallbacks when the page cannot reach the web fonts. An operator who has never seen the system can read the claim and the harm, act on the gate, and open the evidence only if they want to check it.

17 Failure Cases and Limitations

17.1 Known Failure Cases

A fault under the statistical floors is reported as a miss, not as a clean bill of health. A global model or provider fault is invisible to the detectors because the controls stay inside a tenant. Efficiency is measured on resolved sessions, so a fault that changes which conversations resolve would bias it. Cost is read at session grain, so a schema that moved cost onto steps would need a roll-up. CSAT is present on a small and biased subset and is deliberately not surfaced.

17.2 Scale Limitations

At around 100 times the current size the in-memory session table becomes the first limit. The report file is a single-file store that suits one operator, not many concurrent approvers. The window scan grows with the square of the number of days, and the replay service has a run budget. Human review, rather than computation, becomes the binding constraint once the queue is long.

17.3 Data Integrity and Catalog Drift

The catalog disagreed with the corpus in three places and we treated the corpus as the truth for measurement: cost and knowledge-base coverage were declared lower than they are, the calibration set was declared as 100 sessions when 96 join, and the only real emission hole is tool calls. Each affected metric carries a structured declared-versus-observed note, so the disagreement is visible rather than smoothed over. One data blemish, a pasted path inside a label session id, was repaired; the pipeline never reads that file.

Table 8: Proposed fixes.

ID	Finding	Proposed Fix	Predicted Change
P-01	F-02	Validate the <code>get_order_status</code> response so an empty 200 is treated as a failure and takes the existing fallback path	Resolution 0.7755 to 0.8468
P-02	F-01	Publish knowledge-base content for <code>premium_card_info</code> so retrieval returns a confident document	Resolution 0.3087 to 0.6674
P-03	F-03	Edit the <code>acme_main_v3</code> prompt so it confirms once rather than on every turn, keeping the safety wording	Median turns 6.0 to 4.2

Table 12: Threats to inference and how we guard against them.

Threat	Why it matters	How we guard
Multiple comparisons	The scan maximises over roughly 1,500 candidate windows per series and dozens of cohorts, so an unadjusted z of 4 is much weaker as a family-wise test than it appears	Absolute-lift floors plus high z bars
Dependence between sessions	Sessions cluster by day, agent and version, so treating them as independent makes the z optimistic	We lean on effect size rather than the p -value
Winner’s curse	The window is chosen because it contrasts most, so the effect measured on it is inflated	The frozen golden set removes hindsight for fixes; detection stays exposed
Degenerate baselines	An outside rate of exactly zero makes the z a function of a zero-variance baseline rather than of precision	A very large z there is read as directional, not as a probability
Two variance estimators	Detection uses an unpooled standard error and the mix guard a pooled one, so one nominal threshold is not one test	Both thresholds are stated where they apply
Estimand	Efficiency is detected on mean turns and reported as a median, and heavy tails make them disagree	Both are printed
Confidence values	The confidence on a diagnosis is an author prior and does not move with the sample size	Read as an ordering, not a probability
Selection on resolution	Efficiency is measured only on sessions that resolved, so a fault that changes which conversations resolve would bias it	Stated as a limit; the current fault keeps resolution flat
Controls can move with the fault	A tenant-wide KB problem or a change that moves every sibling would defeat the controls, and we would under-report	Accepted trade; we prefer a miss to a false alarm
No cross-tenant control	Peers and siblings stay inside a tenant, which avoids tenant confounds but hides a global model fault	Accepted trade; only the judge-boundary check crosses tenants

17.4 Defects Found and Fixed During Development

We audited the pipeline against a set of edge cases and found real defects, which we record because the fixes are part of the engineering result.

None of these changed the machine score, which is the point: the score alone would not have surfaced them.

18 Assumptions and Transparency

- Everything is implemented in the pipeline except the replay service, which is provided and is a seeded hypothesis verifier.
- No language model is called. Logged facts stay in deterministic Python.
- The judged quality reading comes from the corpus’s own judge and is calibrated against 96 human labels.
- Nothing is simulated inside the pipeline. The replay’s before figure is the endpoint’s own baseline and is labelled as such.
- The final action is human. The system proposes and verifies; a person approves, rejects or defers with a reason.
- No metric definition, judge version or reviewed good set is changed to improve the reported result.

19 Reproducibility

19.1 Execution Procedure

1. **Environment:** Python 3.9 or later; no third-party packages; the provided kit.
2. **Dataset and configuration:** the kit at `kit/`; start the replay service with `tools/nexus-loop-kit/replay/serve.py`.
3. **System execution:** `python3 run.py --kit kit --team NAME --out out/loop-report.json`.
4. **Report generation:** the pipeline writes and validates `out/loop-report.json` against the provided schema.
5. **Validation:** the pipeline scores the report when ground truth is present.
6. **Screen:** `python3 -m nexus_loop.app --report out/loop-report.json`.

19.2 Generated Artifacts

The submission contains the generated `loop-report.json`, the complete system source code, the operator-facing screen, the frozen golden set, the decision log, this report, and a README with the same steps. The final report is generated by the system without manual edits.

Conclusion

Nexus Loop connects detection, diagnosis, prescription, verification and a human decision into one deterministic loop. On the practice corpus it finds three genuine regressions with the right slice and cause, dismisses three lookalikes with written reasons, refuses two questions the data cannot answer, and verifies a fix through replay. It scores 55.0 of 55 machine points and the same on every held-out corpus we generated, with no corpus literal in the detection code.

What makes the loop useful is the provenance behind each number, the evidence behind each diagnosis, verification before a fix is proposed for shipping, a specification instead of a guess when the data cannot answer a question, and a final decision that stays with a person.

References

- [1] Yellow.ai, *Nexus Loop*, *TECHQUEST Problem Statement Series*, Product Development and Agentic AI.
- [2] Yellow.ai, *Nexus Loop Kit*, *Schema*, *Examples and Replay Service*, Provided Challenge Resources.




Figure 5: Summary of final detection, diagnosis and verification results.




Figure 6: Final operator-facing screen showing the finding, evidence, proposed fix, and approval/rejection control.

Table 16: Defects found and fixed during development.

Defect	Effect	Fix
A loop variable shadowed the all-session list	Judged quality, the failover proxy count and the cardinality count silently covered one tenant only	Renamed the loop variable; counts now cover both tenants
The mix-adjusted month comparison mixed two intent scopes	The change beyond mix partly measured a change in which intents were included	Both ends now sit on the common-intent scope; a cohort born in the second month is named
Missing cost was labelled as infinite unit cost	An intent with a telemetry gap read as the worst case	Separated unbounded, which means zero resolutions, from cost unavailable
Empty tool success treated a missing field as zero	A valid response with no result-field key would be flagged as a failure	Only an explicit count of zero counts as empty
The mix dismissal compared two different outcomes	A by-design handoff shift could be dismissed as a mix shift	Both sides of the comparison use the same outcome
The plan advertised breakdowns it did not produce	A reader looking for an advertised cut would not find it	Every declaration was narrowed to what is produced
A judged metric could vanish with no explanation	A missing calibration removed the reading silently	It now files an explicit gap
Regeneration reset recorded approvals	A human verdict was lost on the next run	Approvals are carried forward by identity
Two decisions at once could collide on one temp file	Both returned success while one was lost, and a partial write left the report as invalid JSON	Decisions serialize on a lock; every writer uses a unique temp file, fsync and replace
A corrupt report dropped the connection	The screen could not load and never recovered	The report endpoint returns 500 report_invalid
A non-JSON replay response crashed the run	An unexpected body from the replay endpoint aborted the run before the report was written	The client treats an unreadable body as unavailable
One label session id contained a pasted path	A visibly broken data file, though the pipeline never reads it	The id was restored against the corpus